

SYCL I

Soner Steiner

Intel certified oneAPI Instructor
VSC/TU-WIEN

Overview

- Introduction
- Remainder of the lambda functions
- Compilation and run
- Queues and device selectors
- Manage the data transfer
 - Buffers and Unified Shared Memory
- Basic parallel kernels
- ND-Range kernels
- Sub-groups

What is SYCL?

What is oneAPI Implementation of SYCL?

oneAPI Implementation of SYCL = C++ and SYCL* standard and extensions

Based on modern C++

- C++ productivity benefits and familiar constructs

Standards-based, cross-architecture

- Incorporates the SYCL standard for data parallelism and heterogeneous programming

Extends SYCL* standard

Enhance Productivity

- Simple things should be simple to express
- Reduce verbosity and programmer burden

Enhance Performance

- Give programmers control over program execution
- Enable hardware-specific features

Fast-moving open collaboration feeding into the SYCL* standard

- Open source implementation with goal of upstream LLVM
- Extensions aim to become core SYCL*, or Khronos* extensions

Why not CUDA?

- Unlike CUDA, SYCL supports data parallelism in C++ for all vendors and all types of architectures (not just GPUs).
- CUDA is focused on NVIDIA GPU support only, and efforts (such as HIP/ROCm) to reuse it for GPUs by other vendors have limited ability.
- With the explosion of accelerator architectures, only SYCL offers the support we need for harnessing this diversity and offering a multivendor/multiarchitecture approach to help with portability that CUDA does not offer.

new-golden-age-for-computer-architecture

Why Standard C++ with SYCL?

- Every program using SYCL is first and foremost a C++ program.
- SYCL takes C++ programming places it cannot go without SYCL.
- We don't believe the C++ standard will evolve to displace the need for SYCL anytime soon!?

Getting a C++ compiler with SYCL support?

- Compilers supporting SYCL: **khronos-sycl**
- By using LLVM, the DPC++ compiler project has backends for numerous devices.
- This has already resulted in support for Intel, Nvidia and AMD GPUs, numerous CPUs and Intel FPGAs.
- oneAPI Tools, including the libraries, debuggers, DPC++ compiler and other tools, which are freely available.

Data Parallel C++

Standards-based, Cross-architecture Language

DPC++

=

ISO C++

+

Khronos SYCL

+

Community Extensions

tinyurl.com/sycl2020-support-in-dpcpp



Direct Programming:
Data Parallel C++

Community Extensions

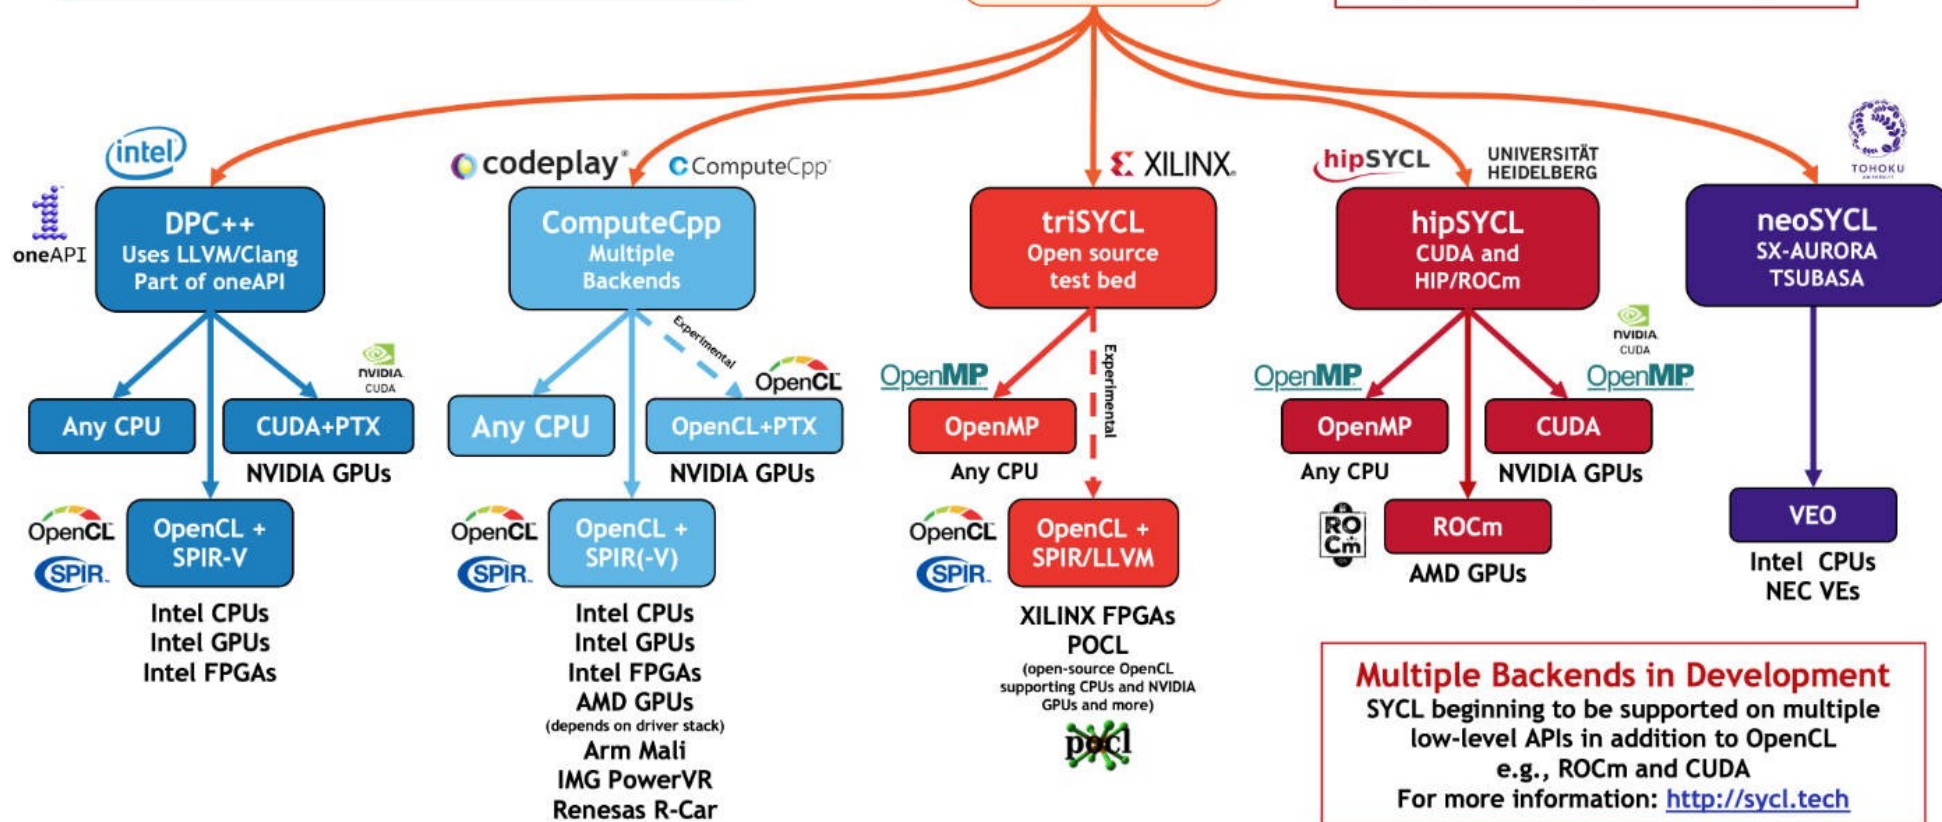
Khronos SYCL

ISO C++

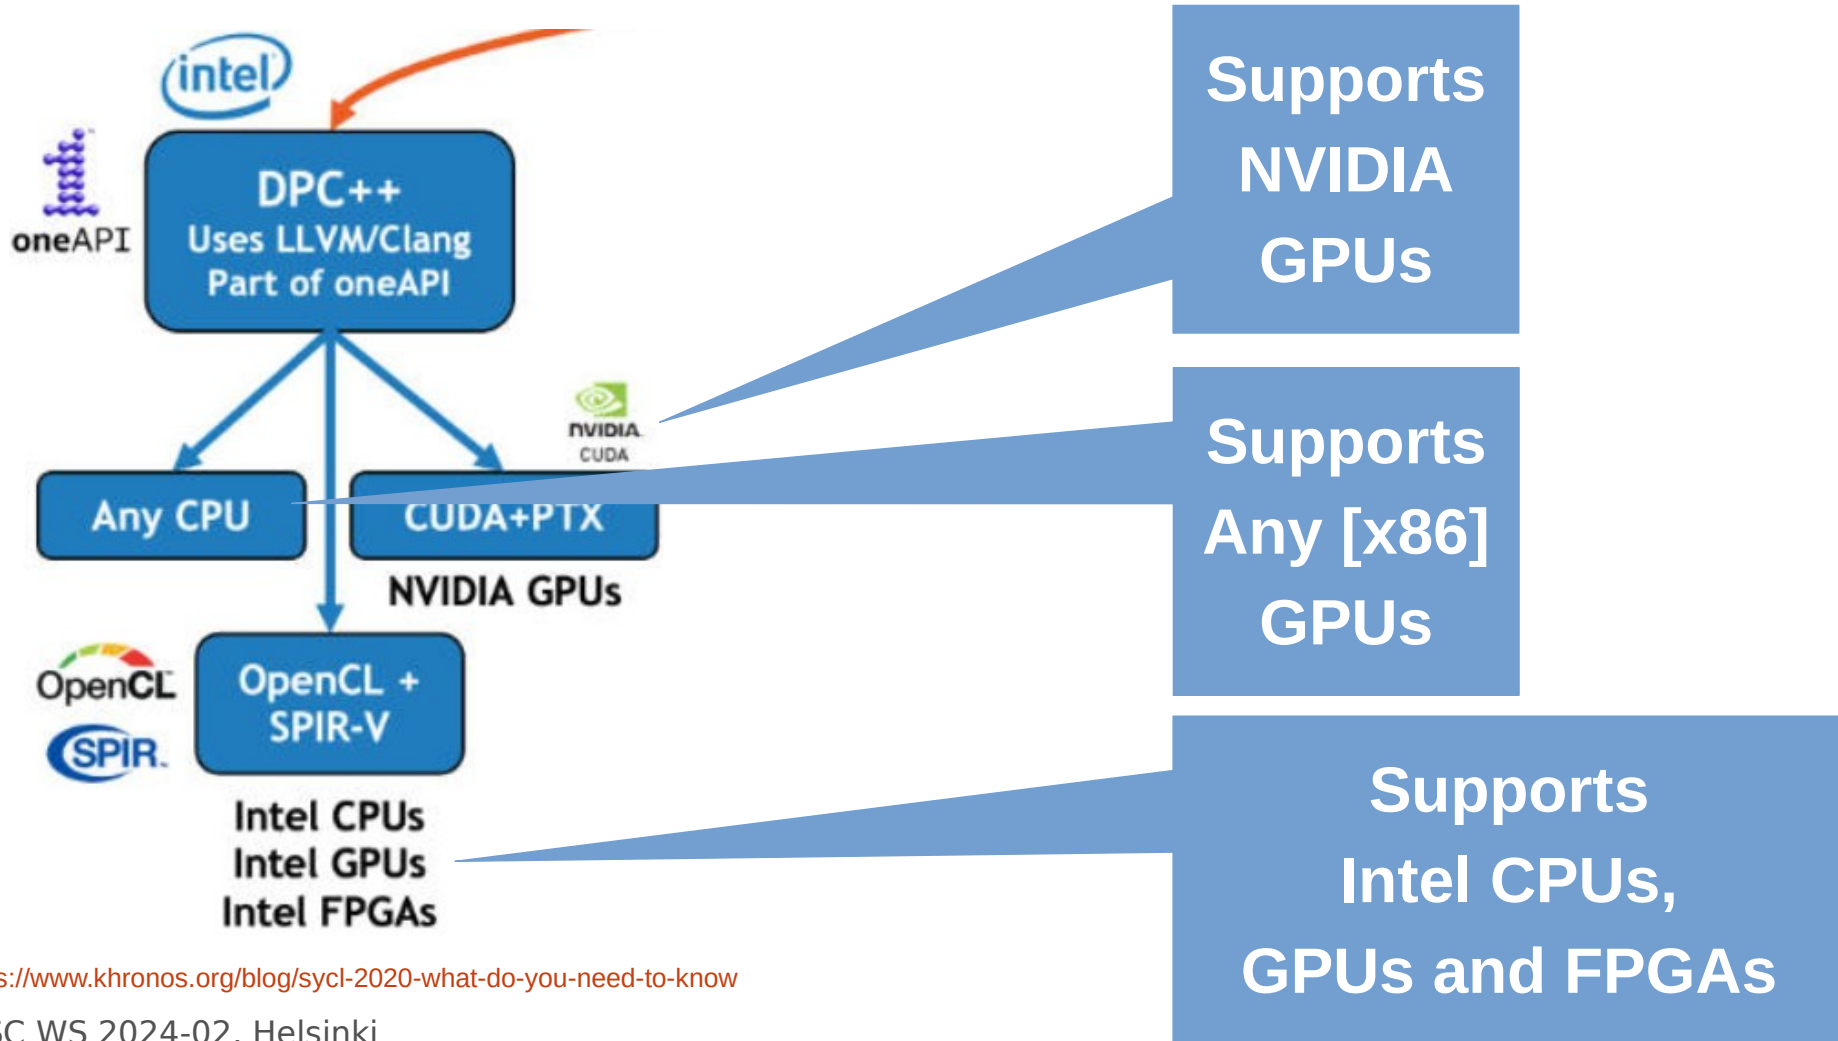
Intel DC++ in the SYCL ecosystem?

SYCL, OpenCL and SPIR-V, as open industry standards, enable flexible integration and deployment of multiple acceleration technologies

SYCL enables Khronos to influence ISO C++ to (eventually) support heterogeneous compute



Intel DC++ in the SYCL ecosystem?




Many of the source examples are from Book:

- Source code accessible from
\$ oneapi-cli

```
(1) Create a project
(2) View oneAPI docs in browser
(q) Quit
```

```
(1) cpp
(2) python
(3) fortran
(b) Back
(q) Quit
```



data-parallel-c++-free-book-2ed

Samples	Description
<ul style="list-style-type: none"> Compiler Infrastructure <ul style="list-style-type: none"> Intrinsics Matrix Multiply OpenMP Offload Graph Traversal <ul style="list-style-type: none"> MergeSort OMP Parallel Patterns <ul style="list-style-type: none"> OpenMP* Reduction Structured Grids <ul style="list-style-type: none"> ISO3DFD OMP Offload Tutorials Jupyter Notebooks <ul style="list-style-type: none"> OpenMP Offload C++ Tutorials C++SYCL <ul style="list-style-type: none"> Combinational Logic <ul style="list-style-type: none"> Mandelbrot Sepia Filter Concurrent Kernels convolutionSeparable Dense Linear Algebra <ul style="list-style-type: none"> Base: Vector Add Complex Mult Jacobi Cuda Graphs Jacobi Iterative Solver Matrix Multiply 	<p>This sample Multiplies two large Matrices in parallel using SYCL and OpenMP* (OMP)</p> <p>The following tools are needed to build this sample but are not locally installed: (icc)</p> <p>You may continue and view the sample without the prerequisites. To install the missing prerequisites, visit:</p> <p>https://www.intel.com/content/www/us/en/developer/tools/oneapi/overview-hpc-kit</p> <p>or https://www.intel.com/content/www/us/en/developer/tools/oneapi/overview-wiot-kit</p> <p>Press Backspace to return to previous screen!</p>

Compilation and Run

prompt

```
$ source /opt/intel/oneapi/setvars.sh
```

```
$ dpcpp -O2 -g -std=c++17 -o 00Hello.out 00Hello.cpp
```

!

NOW

```
$ icpx -fsycl -O2 -g -std=c++17 -o 00-Hello.x 00-Hello.cpp
```

List SYCL Devices available

```
$ sycl-ls [--verbose]
```

Where Code Executes

```
$sycl-ls  
[opencl:acc:0] Intel(R) FPGA Emulation Platform for OpenCL(TM), Intel(R) FPGA  
Emulation Device 1.2 [2022.15.12.0.01_081451]  
[opencl:cpu:1] Intel(R) OpenCL, Intel(R) Core(TM) i9-10900K CPU @ 3.70GHz 3.0  
[2022.15.12.0.01_081451]  
[opencl:gpu:2] Intel(R) OpenCL HD Graphics, Intel(R) UHD Graphics 630 [0x9bc5]  
3.0 [21.36.20889]
```

three devices on this computer

Where Code Executes

```
$:> sycl-ls  
[opencl:cpu:0] Intel(R) OpenCL, 11th Gen Intel(R) Core(TM) i7-11800H @ 2.30GHz OpenCL 3.0 (Build 0) [2023.16.10.0.17_160000]  
[opencl:acc:1] Intel(R) FPGA Emulation Platform for OpenCL(TM), Intel(R) FPGA Emulation Device OpenCL 1.2 [2023.16.6.0.22_223734]  
[opencl:cpu:2] Intel(R) OpenCL, 11th Gen Intel(R) Core(TM) i7-11800H @ 2.30GHz OpenCL 3.0 (Build 0) [2023.16.6.0.22_223734]  
[ext_oneapi_cuda:gpu:0] NVIDIA CUDA BACKEND, NVIDIA GeForce RTX 3070 Laptop GPU 8.6 [CUDA 11.4]
```

four devices on my computer

Where Code Executes

```
$:>  
$:> ssh -i lumi-rsa steiners2@mahti.csc.fi  
steiners2@mahti.csc.fi's password:  
Welcome _____  
CSC - Tieteen tietotekniikan keskus - IT Center for Science  
  
  MAHTI  
  
Mahti.csc.fi - Atos BullSequana XH2000  
1404 AMD Rome CPU nodes - 24 Nvidia A100 GPU nodes  
Contact _____
```

Where Code Executes

```
[steiners2@mahti-login15 ~]$ . /scratch/project_2008874/cristian/intel/oneapi/setvars.sh --include-intel-llvm
:: initializing oneAPI environment ...
-bash: BASH_VERSION = 4.4.20(1)-release
args: Using "$@" for setvars.sh arguments: --include-intel-llvm
:: advisor -- latest
:: ccl -- latest
:: compiler -- latest
:: dal -- latest
:: debugger -- latest
:: dev-utilities -- latest
:: dnnl -- latest
```

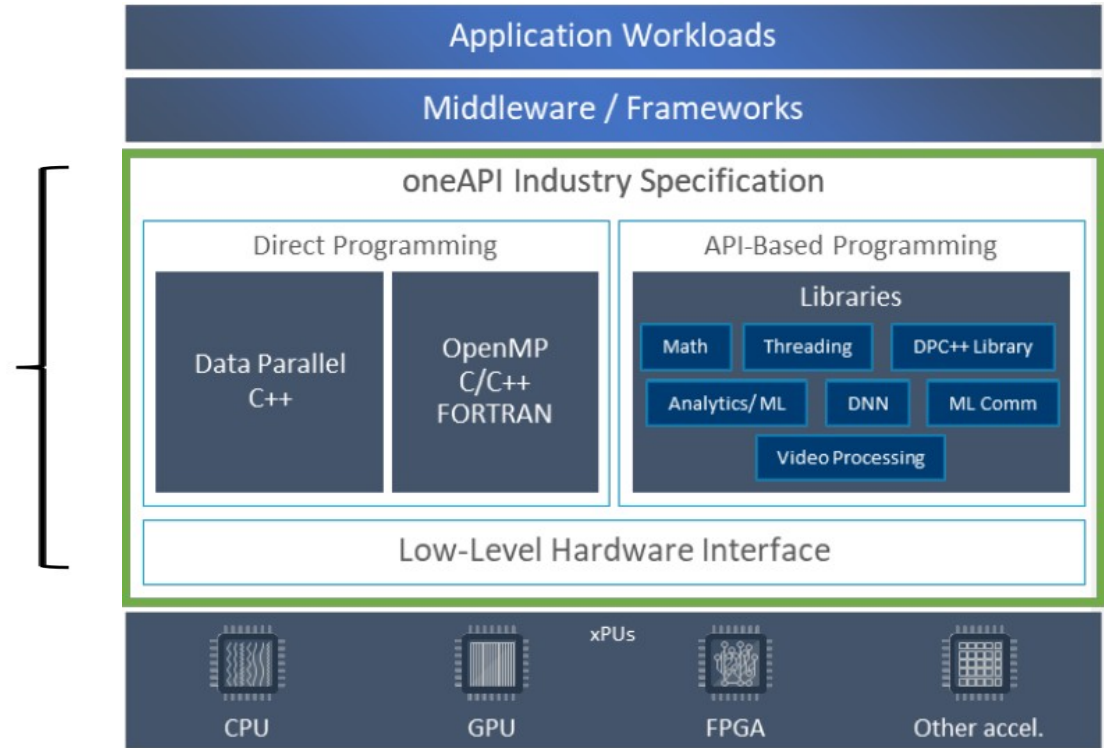
```
[steiners2@mahti-login15 ~]$ module load cuda
[steiners2@mahti-login15 ~]$ sycl-ls
[opencl:acc:0] Intel(R) FPGA Emulation Platform for OpenCL(TM), Intel(R) FPGA Emulation Device OpenCL 1.2 [2023.16.12.0.12_195853.xmain-hotfix]
[opencl:cpu:1] Intel(R) OpenCL, AMD EPYC 7402 24-Core Processor OpenCL 3.0 (Build 0) [2023.16.12.0.12_195853.xmain-hotfix]
```

Where Code Executes

[illegible]

Programmers' perspective: Three things to consider

1. Offload the code to device
2. Manage the transfer of Data
3. Implement Parallelism

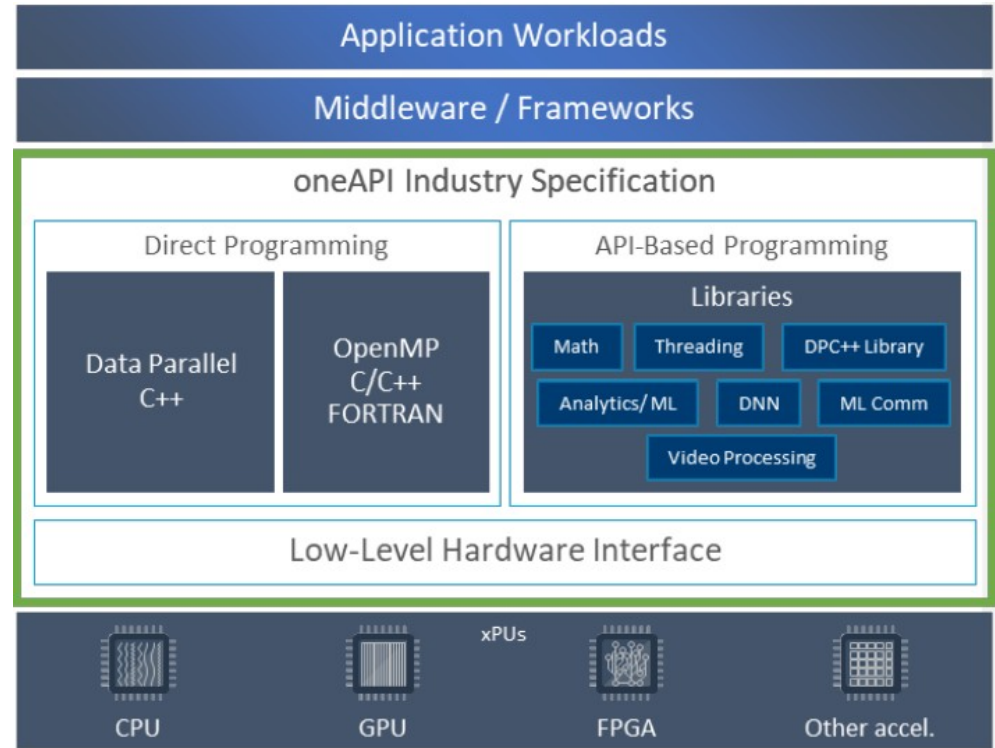


Programmers' perspective: Three things to consider

1. Offload the code to
device

2. Manage the transfer
of Data

3. Implement
Parallelism



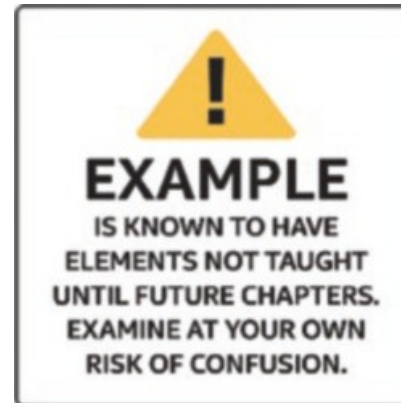
```

1 #include<iostream>
2 #include<sycl/sycl.hpp>
3 using namespace sycl;
4
5 const std::string secret
6 {
7     "Ifmmp-!xpsme\""\012J(n!tpssz-!Ebwf/!"
8     "J(n!bgsbje!J!dbo(u!ep!uibu/!..IBM\01"
9 };
10
11 const auto sz=secret.size();
12
13 int main()
14 {
15     queue Q;
16     char* result = malloc_shared<char>(sz, Q);
17     std::memcpy(result, secret.data(), sz);
18
19     Q.parallel_for(sz, [=] (auto& i)
20     {
21         result[i] -= 1;
22     }).wait();
23
24     std::cout << result << "\n" ;
25     free(result, Q);
26     return 0;
27 }

```

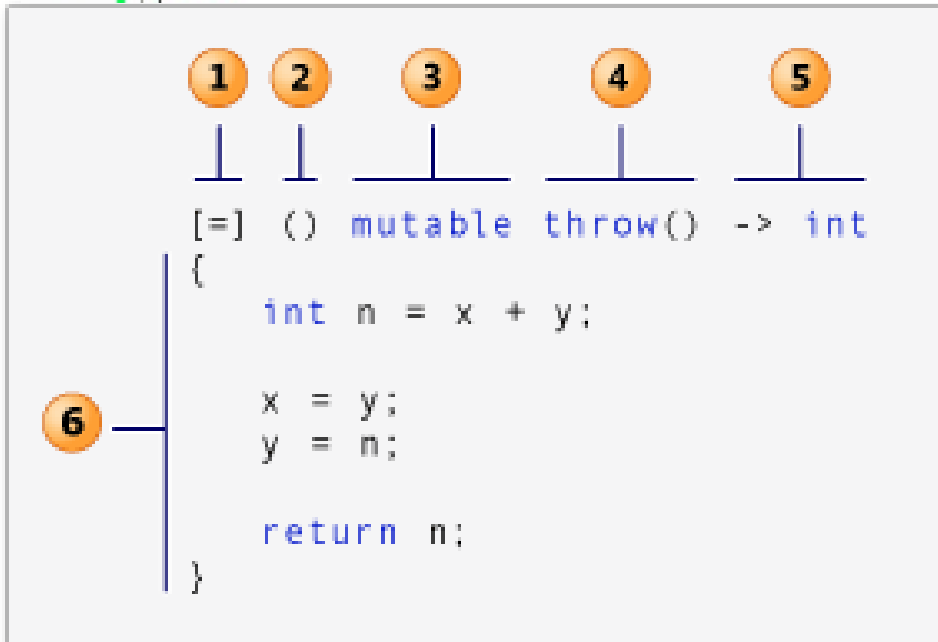
SYCLs Hello World

- 1: Access to all SYCL constructs
- 3: Avoid having to write sycl::
- 15: Establish queue for work
requests to a particular device
- 16: create shared data
- 19: Enqueue work to the device
- 21: Only line that runs on the device



Lambda-functions ... Lambdas

```
39 q.parallel_for(N, [=](auto i)
40 {
41     a[i] -= 2;
42 });
```



1. **capture clause**
 2. **parameter list** optional
 3. **mutable**
specification optional
 4. **exception-**
specification optional
 5. **trailing-return-**
type optional
 6. **lambda body**
- `[=]` : capture by value
 - `[&]` : capture by reference

<https://learn.microsoft.com/en-us/cpp/cpp/lambda-expressions-in-cpp>

Kernel Code

```
39 q.parallel_for(N, [=](auto i)
40 {
41     a[i] -= 2;
42 });
```

Kernel Code
Cannot use
these features



- Run Asynchronously
- Limitation on what kind of C++ code
 - Dynamic Polymorphism
 - Dynamic memory allocations
 - Static variables
 - Function pointers
 - Runtime Type Information (RTTI)
 - Exception Handling
 - Recursion

SYCL fundamentals

Explain the SYCL fundamental classes

Use **device selection** to **offload kernel workloads**

Decide when to use basic parallel kernels and ND-Range kernels

Understand various ways to **synchronize** data between host and device with using buffer memory model

Write a complete SYCL program that offload computation to accelerator device

C++ with SYCL

- Enables programming for **heterogenous hardware** from **different vendors**.
- **Single source** that has host code and kernel code to offload to CPU, GPU, FPGA or other accelerator devices.
- Based on Open Standards C++ and Khronos* SYCL

Anatomy of a SYCL Application

```
#include <sycl/sycl.hpp>
using namespace sycl;

int main() {
    std::vector<float> A(1024, 1.0f), B(1024, 2.0f), C(1024);
    {
        buffer bufA {A}, bufB {B}, bufC {C};
        queue q;
        q.submit([&](handler &h) {
            auto A = bufA.get_access(h, read_only);
            auto B = bufB.get_access(h, read_only);
            auto C = bufC.get_access(h, write_only);
            h.parallel_for(1024, [=](auto i){
                C[i] = A[i] + B[i];
            });
        });
    }
    for (int i = 0; i < 1024; i++)
        std::cout << "C[" << i << "] = " << C[i] << std::endl;
}
```

Host code

Accelerator
device code

Host code

Anatomy of a SYCL Application

```
#include <sycl/sycl.hpp>
using namespace sycl;

int main() {
    std::vector<float> A(1024, 1.0f), B(1024, 2.0f), C(1024);
    {
        buffer bufA {A}, bufB {B}, bufC {C};
        queue q;
        q.submit([&](handler &h) {
            auto A = bufA.get_access(h, read_only);
            auto B = bufB.get_access(h, read_only);
            auto C = bufC.get_access(h, write_only);
            h.parallel_for(1024, [=](auto i){
                C[i] = A[i] + B[i];
            });
        });
    }
    for (int i = 0; i < 1024; i++)
        std::cout << "C[" << i << "] = " << C[i] << std::endl;
}
```

Application scope

Command group
scope

Device scope

Application scope

SYCL Basics

```
std::vector<float> A(1024, 1.0f), B(1024, 2.0f), C(1024);
{
    buffer bufA {A}, bufB {B}, bufC {C};
    queue q;
    q.submit([&](handler &h)
    {
        auto A = bufA.get_access(h, read_only);
        auto B = bufB.get_access(h, read_only);
        auto C = bufC.get_access(h, write_only);
        h.parallel_for(1024, [=](auto i)
        {
            C[i] = A[i] + B[i];
        });
    });
}

for (int i = 0; i < 1024; i++)
    std::cout << "C[" << i << "] = " << C[i] << std::endl;
}
```

Buffers creation via
host vectors/pointers

Buffers encapsulate
data
in a SYCL application

- Across both devices
and host!

SYCL Basics

```
std::vector<float> A(1024, 1.0f), B(1024, 2.0f), C(1024);
```

```
{  
    buffer bufA {A}, bufB {B}, bufC {C};  
    queue q;  
    q.submit([&](handler &h)  
    {  
        auto A = bufA.get_access(h, read_only);  
        auto B = bufB.get_access(h, read_only);  
        auto C = bufC.get_access(h, write_only);  
        h.parallel_for(1024, [=](auto i)  
        {  
            C[i] = A[i] + B[i];  
        }));  
    });  
};
```

```
for (int i = 0; i < 1024; i++)  
    std::cout << "C[" << i << "] = " << C[i] << std::endl;  
}
```

- A queue submits command groups to be executed by the SYCL runtime
- Queue is a mechanism where work is submitted to a device.

SYCL CLASSES

Where Code Executes

- Queues
- Device Selectors

QUEUES CONNECT US TO DEVICES

- We submit actions into queues to request computational work and data movement.
- Actions happen **ASYNCHRONOUSLY**

Device

- The **device** class represents the capabilities of the accelerators in a oneAPI system.
- The device class contains member functions for **querying information about the device**, which is useful for DPC++ programs where multiple devices are created.
- The function **get_info** gives information about the device:
 - Name, vendor, and version of the device
 - The local and global work item IDs
 - Width for built in types, clock frequency, cache width and sizes, online or offline

```
queue q;  
device my_device = q.get_device();  
std::cout << "Device: " << my_device.get_info<info::device::name>() << std::endl;
```

Device Selector

- The **device_selector** class enables the runtime selection of a particular device to execute kernels based upon user-provided heuristics.
- The following code sample shows use of the standard device selectors (**default_selector**, **cpu_selector**, **gpu_selector**...) and a derived **device_selector**

```
default_selector_v selector;  
// host_selector_v selector;  
// cpu_selector_v selector;  
// gpu_selector_v selector;  
queue q(selector);  
std::cout << "Device: " << q.get_device().get_info<info::device::name>() << std::endl;
```

Queue

- A queue **submits command groups** to be executed by the SYCL runtime
- Queue is a mechanism where work is submitted to a device.
- A Queue map to one device and multiple queues can be mapped to the same device.

```
queue q;
```

```
q.submit( [&] (handler& h)  
{  
    // COMMAND GROUP CODE  
});
```

The queue class

Actions are submitted to a queue for execution on a single device

- Always bound to a single device
 - Q1 → GPU1
 - Q2 → CPU
 - Q3 → FPGA
 - Q4 → GPU
- Several queues can point to the same device
 - Q1 → GPU1
 - Q2 → GPU1
 - Q3 → GPU1
 - Q4 → CPU

Choosing Where Device Kernels Run

Work is submitted to queues

- Each queue is associated with exactly one device (e.g. a specific GPU or FPGA)
- You can:
 - Decide which device a queue is associated with (if you want)
 - Have as many queues as desired for dispatching work in heterogeneous systems

Create queue targeting any device:	<pre>queue();</pre>
Create queue targeting a pre-configured classes of devices:	<pre>queue(cpu_selector_v); queue(gpu_selector_v); queue(ext::intel::fpga_selector_v); queue(accelerator_selector_v); queue(host_selector_v);</pre>
Create queue targeting specific device (custom criteria):	<pre>class custom_selector : public device_selector { int operator()(..... // Any logic you want! ... queue(custom_selector);</pre>

Always
available



The queue class – Binding done at construction

```
class queue {
public:
    // Create a queue associated with a default
    // (implementation chosen) device.
    queue(const property_list & = {});

    queue(const async_handler &, const property_list & = {});

    // Create a queue using a DeviceSelector.
    // A DeviceSelector is a callable that ranks
    // devices numerically. There are a few SYCL-defined
    // device selectors available such as
    // cpu_selector_v and gpu_selector_v.
    template <typename DeviceSelector>
    explicit queue(const DeviceSelector &deviceSelector,
                  const property_list &propList = {});

    // Create a queue associated with an explicit device to
    // which the program already holds a reference.
    queue(const device &, const property_list & = {});

    // Create a queue associated with a device in a specific
    // SYCL context. A device selector may be used in place
    // of a device.
    queue(const context &, const device &,
          const property_list & = {});
};
```

**Default
Device
used here**

The queue class – key member functions

a queue is
bound to a
single devices

```
class queue {  
public:  
    // Submit a command group to this queue.  
    // The command group may be a lambda expression or  
    // function object. Returns an event reflecting the status  
    // of the action performed in the command group.  
    template <typename T>  
    event submit(T);  
  
    // Wait for all previously submitted actions to finish  
    // executing.  
    void wait();  
  
    // Wait for all previously submitted actions to finish  
    // executing. Pass asynchronous exceptions to an  
    // async_handler function.  
    void wait_and_throw();  
};
```


Device selectors

Selector
cpu_selector_v
gpu_selector_v
ext::intel::fpga_selector_v
accelerator_selector_v
default_selector_v
<i>or write a custom selector</i>

Choosing Devices: Five use cases:

#	Methods	Comments
1	Anywhere (don't care where)	Runtime chooses
2	Always on Host	Good for debugging
3	GPU or Accelerator	
4	Heterogeneous set of devices	
5	Specific Class of device	e.g. FPGA

Method #1, Binding a Queue to a Device When Any Device Will Do

Default Device
used here
Decided by the
runtime

```
1 #include <sycl/sycl.hpp>
2 #include <iostream>
3 using namespace sycl;
4
5 int main()
6 {
7     // create queue on whatever default device that the
8     // implementation chooses.
9     queue Q;
10
11     std::cout << "Selected device: " <<
12     Q.get_device().get_info<info::device::name>() << "\n";
13
14     return 0;
15 }
16
```



Method #1, Binding a Queue to a Device When Any Device Will Do

Sample Outputs (one line per run depending on system):

Selected device: NVIDIA GeForce RTX 3060

Selected device: AMD Radeon RX 5700 XT

Selected device: Intel(R) Data Center GPU Max 1100

Selected device: Intel(R) FPGA Emulation Device

Selected device: AMD Ryzen 5 3600 6-Core Processor

Selected device: Intel(R) UHD Graphics 770

Selected device: Intel(R) Xeon(R) Gold 6128 CPU @ 3.40GHz

Selected device: 11th Gen Intel(R) Core(TM) i9-11900KB @ 3.30GHz

many more possible... these are only examples



Method #2 Using Host Device, Development, Debugging and Deployment

```
1 #include <sycl/sycl.hpp>
2 #include <iostream>
3 using namespace sycl;
4
5 int main()
6 {
7     queue Q{ cpu_selector{} };
8
9     std::cout << "Selected Device: " <<
10     Q.get_device().get_info<info::device::name>() << "\n";
11     std::cout << " ---->>>> Device Vendor: " <<
12     Q.get_device().get_info<info::device::vendor>() << "\n";
13
14     return 0;
15 }
```

Old notation
SYCL 2020
easier

Method #3 -Using a GPU or Accelerator (just change the selector to **gpu_selector** or **accelerator_selector**)



Method #2 Using Host Device, Development, Debugging and Deployment

SYCL 2020
easier

```
1 #include <sycl/sycl.hpp>
2 #include <iostream>
3 using namespace sycl;
4
5 int main()
6 {
7     queue Q{ cpu_selector_v};
8     // queue Q2{ host_selector_v};
9     queue Q3{ default_selector_v};
10    queue Q4{ gpu_selector_v};
11
12    std::cout << "Selected Device: " <<
13    Q.get_device().get_info<info::device::name>() << "\n";
14    std::cout << " ---->>>> Device Vendor: " <<
15    Q.get_device().get_info<info::device::vendor>() << "\n";
16
17    std::cout << "Selected Device: " <<
18    Q3.get_device().get_info<info::device::name>() << "\n";
19    std::cout << " ---->>>> Device Vendor: " <<
20    Q3.get_device().get_info<info::device::vendor>() << "\n";
21
22    std::cout << "Selected Device: " <<
23    Q4.get_device().get_info<info::device::name>() << "\n";
24    std::cout << " ---->>>> Device Vendor: " <<
25    Q4.get_device().get_info<info::device::vendor>() << "\n";
26
```



Method #3 -
Using a GPU or Accelerator
just change the selector to
gpu_selector_v
or
accelerator_selector_v



Method #4 Using Multiple Devices

```
1 #include <sycl/sycl.hpp>
2 #include <sycl/ext/intel/fpga_extensions.hpp>
3 #include <iostream>
4 using namespace sycl;
5
6 int main()
7 {
8     queue gpu_q( gpu_selector_v );
9     queue cpu_q( cpu_selector_v );
10    queue fpga_q( ext::intel::fpga_selector_v );
11
12    std::cout << "Selected Device1: " <<
13    cpu_q.get_device().get_info<info::device::name>() << "\n";
14    std::cout << "Selected Device2: " <<
15    gpu_q.get_device().get_info<info::device::name>() << "\n";
16    std::cout << "Selected Device3: " <<
17    fpga_q.get_device().get_info<info::device::name>() << "\n";
18
19    return 0;
20 }
```

Three Queues
Three Devices



Method #5: Custom (Very Specific)
Device Selection → Skip

Control Device Selection via SYCL_DEVICE_FILTER

- Limits the choice of devices available to the runtime
- Syntax: `SYCL_DEVICE_FILTER=backend:device_type:device_num, ...`
 - **Backend**: host, opencl, level_zero, cuda, hip, *
 - **Device_type**: host, cpu, gpu, acc, *
 - **Device_num**: unsigned integer
 - Enumeration index of devices from the sycl-ls utility
 - Each field is *optional*, so missing entry is regarded as '*'.
 - E.g., `SYCL_DEVICE_FILTER=gpu` `SYCL_DEVICE_FILTER=*:gpu:*`
 - Multiple triples can be specified separated by commas.



Control Device Selection via SYCL_DEVICE_FILTER

■ Dual purposes

- Users can specify their desired devices with the given triple(s).
- SYCL only loads relevant plugins into runtime.

Control Device Selection via SYCL_DEVICE_FILTER

```
$ icpx -fsycl 02-Default-selector.cpp -o 02-Default-selector.x
```

```
$ SYCL_PI_TRACE=1 ./02-Default-selector.x
```

```
$ SYCL_PI_TRACE=1 ./02-Default-selector.x
SYCL_PI_TRACE[basic]: Plugin found and successfully loaded: libpi_opencl.so [ PluginVersion: 14.37.1 ]
SYCL_PI_TRACE[basic]: Plugin found and successfully loaded: libpi_cuda.so [ PluginVersion: 14.38.1 ]
SYCL_PI_TRACE[basic]: Plugin found and successfully loaded: libpi_unified_runtime.so [ PluginVersion: 14.37.1 ]
SYCL_PI_TRACE[all]: Requested device_type: info::device_type::automatic
SYCL_PI_TRACE[all]: Requested device_type: info::device_type::automatic
SYCL_PI_TRACE[all]: Requested device_type: info::device_type::automatic
SYCL_PI_TRACE[all]: Selected device: -> final score = 500
SYCL_PI_TRACE[all]:   platform: NVIDIA CUDA BACKEND
SYCL_PI_TRACE[all]:   device: NVIDIA GeForce RTX 3070 Laptop GPU
Selected device: NVIDIA GeForce RTX 3070 Laptop GPU
```



Control Device Selection via SYCL_DEVICE_FILTER

```
$ icpx -fsycl 02-Default-selector.cpp -o 02-Default-selector.x
```

```
$ SYCL_PI_TRACE=1 SYCL_DEVICE_FILTER=*:cpu ./02-Default-selector.x
```

```
SYCL_PI_TRACE[basic]: Plugin found and successfully loaded: libpi_opencl.so [ PluginVersion: 14.37.1 ]  
SYCL_PI_TRACE[basic]: Plugin found and successfully loaded: libpi_cuda.so [ PluginVersion: 14.38.1 ]  
SYCL_PI_TRACE[basic]: Plugin found and successfully loaded: libpi_unified_runtime.so [ PluginVersion: 14.37.1 ]  
SYCL_PI_TRACE[all]: Requested device_type: info::device_type::automatic  
SYCL_PI_TRACE[all]: Selected device: -> final score = 300  
SYCL_PI_TRACE[all]: platform: Intel(R) OpenCL  
SYCL_PI_TRACE[all]: device: 11th Gen Intel(R) Core(TM) i7-11800H @ 2.30GHz  
Selected device: 11th Gen Intel(R) Core(TM) i7-11800H @ 2.30GHz
```



Examples

ONEAPI_DEVICE_SELECTOR

ONEAPI_DEVICE_SELECTOR=

<i>opencl:*</i>	Only the OpenCL devices are available
<i>level_zero:gpu</i>	Only GPU devices on the Level Zero platform are available.
<i>"opencl:gpu;level_zero:gpu"</i>	GPU devices from both Level Zero and OpenCL are available. Note that escaping (like quotation marks) will likely be needed when using semi- colon separated entries.
<i>opencl:gpu,cpu</i>	Only CPU and GPU devices on the OpenCL platform are available.
<i>opencl:0</i>	Only the device with index 0 on the OpenCL backend is available.
<i>hip:0,2</i>	Only devices with indices of 0 and 2 from the HIP backend are available.

Dispatching mechanism

Dispatching Code – Device Dispatch Mechanism

- So far. We've used

- `queue::parallel_for()`
- `queue::single_task()`

```
39  q.parallel_for(N, [=](auto i)
40  {
41      a[i] -= 2;
42  });
```

- `handler::single_task()`
- `handler::parallel_for()`
- `handler::parallel_for_work_group()`

Kernel

- The kernel class encapsulates methods and data for executing code on the device when a command group is instantiated
- Kernel object is not explicitly constructed by the user
- Kernel object is constructed when a kernel dispatch function, such as `parallel_for`, is called

```
q.submit( [&] (handler& h)
{
    h.parallel_for(N, [=](auto i)
    {
        A[i] = B[i] + C[i]);
    });
});
```

Language Simplification

Code snippet below shows how SYCL* code can be simplified

```
buffer<int, 1> buf(data.data(), data.size());  
q.submit([&] (handler &h){  
    auto A = buf.get_access<access::mode::read_write>(h);  
    h.parallel_for<class kernel>(range<1>(N), [=](id<1> i)  
    { A[i] += 1; }));  
});
```

Annotations for SYCL code:

- Buffer Simplification (points to `buffer<int, 1>`)
- Accessor Simplification (points to `buf.get_access`)
- Lambda name no longer required (points to `class kernel`)
- parallel_for simplification (points to `parallel_for`)

SYCL

```
buffer buf(data);  
q.submit([&] (handler &h){  
    auto A = accessor(buf, h);  
    h.parallel_for(N, [=](auto i)  
    { A[i] += 1; }));  
});
```

← *Simple and
Less
Verbose*

SYCL 2020

DPC++ language and runtime

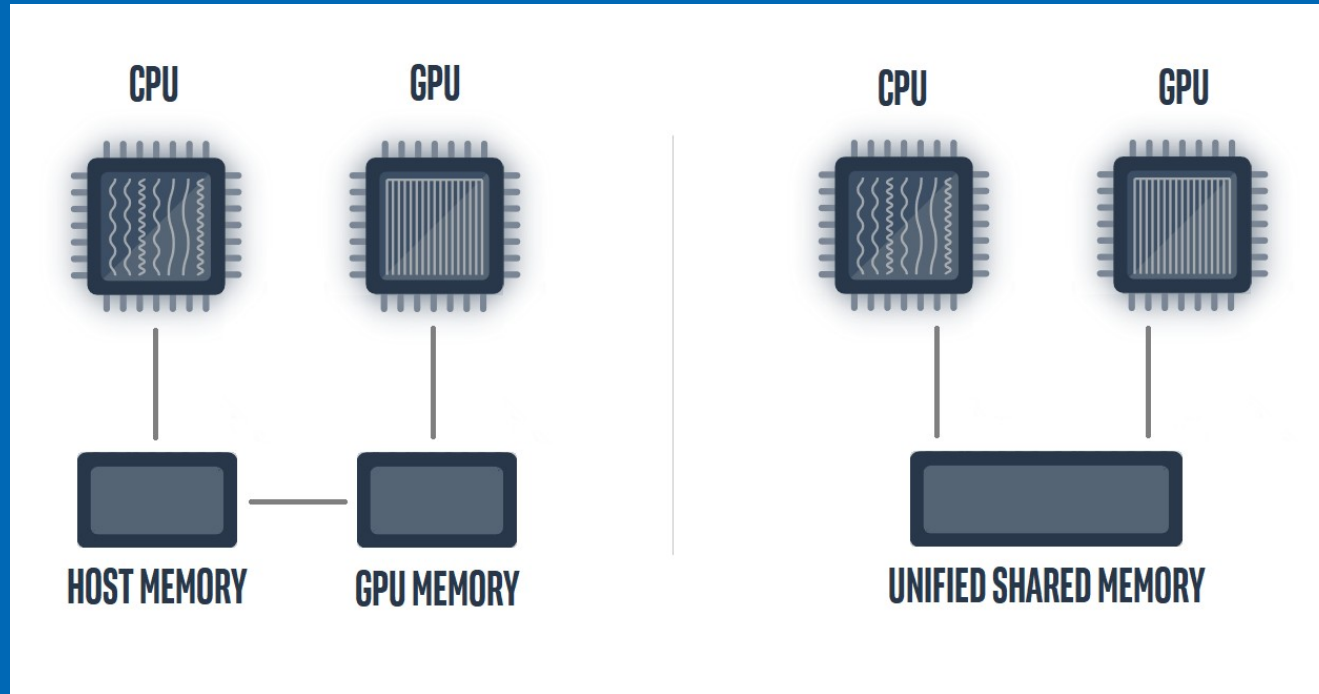
- DPC++ language and runtime consists of a set of C++ classes, templates, and libraries
- **Application scope** and **command group scope** :
 - Code that executes on the host
 - The full capabilities of C++ are available at application and command group scope
- **Kernel scope**:
 - Code that executes on the device.
 - At kernel scope there are limitations in accepted C++

Actions

Work Type	Actions (handler class methods)	Summary
Device code execution	single_task	Execute a single instance of a device function.
	parallel_for	Multiple forms are available to launch device code with different combinations of work sizes.
	parallel_for_work_group	Launch a kernel using hierarchical parallelism, described in Chapter 4.
Explicit memory operation	copy	Copy data between locations specified by accessor, pointer, and/or shared_ptr. The copy occurs as part of the DAG, including dependence tracking.
	update_host	Trigger update of host data backing of a buffer object.
	fill	Initialize data in a buffer to a specified value.

Developer View of USM

Developers can reference **same memory object** in host and device code with Unified Shared Memory



Unified Shared Memory (USM)

Unified Shared Memory can be setup as follows:

```
int *data = malloc_shared<int>(N, q);
```

You can also use a more familiar C++/C style malloc:

```
int *data = static_cast<int*>(malloc_shared(N * sizeof(int), q));
```

Unified Shared Memory

Unified Shared Memory enables accessing memory on the host and device with same pointer reference

Setup Unified
Shared Memory

```
queue q;  
auto data = malloc_shared<int>(N, q);
```

Host can
initialize

```
for(int i=0;i<N;i++) data[i] = 10;
```

Device can
modify

```
q.parallel_for(N, [=](auto i)  
{  
    data[i] += 1;  
}).wait();
```

Host has output

```
for(int i=0;i<N;i++) std::cout << data[i] << " ";  
free(data, q);
```

Exercises

■ SYCL Program Structure

- Read the instructions carefully, it is about compiling and env variables.
- The sycl I exercises can be done on the Intel Dev Cloud (IDC) and or on LUMI.

Getting Started on DevCloud

- `qsub -l -l nodes=1:gpu:ppn=2 -d .`
- `sycl-ls` (control devices via `SYCL_DEVICE_FILTER`)
- Compile and run a simple `vecAdd` code
- `export SYCL_PI_TRACE=1`
- `export ONEAPI_DEVICE_SELECTOR=opencl:gpu`

Useful Links

Open source projects

oneAPI Data Parallel C++ compiler: github.com/intel/llvm

Graphics Compute Runtime: Graphics github.com/intel/compute-runtime

Compiler: github.com/intel/intel-graphics-compiler

SYCL 2020: tinyurl.com/sycl2020-spec

DPC++ Extensions: tinyurl.com/dpcpp-ext

Environment Variables: tinyurl.com/dpcpp-env-vars

DPC++ book: tinyurl.com/dpcpp-book

SYCL Academy github.com/codeplaysoftware/syclacademy/tree/main

Code samples:
github.com/intel/llvm/tree/sycl/sycl/test
github.com/intel/llvm/tree/sycl/sycl/test-e2e
github.com/oneapi-src/oneAPI-samples

Notices and Disclaimers

- Performance varies by use, configuration and other factors. Learn more at www.Intel.com/PerformanceIndex
- Performance results are based on testing as of dates shown in configurations and may not reflect all publicly available updates. See backup for configuration details. No product or component can be absolutely secure.
- You may not use or facilitate the use of this document in connection with any infringement or other legal analysis concerning Intel products described herein. You agree to grant Intel a non-exclusive, royalty-free license to any patent claim thereafter drafted which includes subject matter disclosed herein.
- The products described may contain design defects or errors known as errata which may cause the product to deviate from published specifications. Current characterized errata are available on request.
- No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document, with the sole exception that a) you may publish an unmodified copy and b) code included in this document is licensed subject to the Zero-Clause BSD open source license (0BSD), <https://opensource.org/licenses/0BSD>. You may create software implementations based on this document and in compliance with the foregoing that are intended to execute on the Intel product(s) referenced in this document. No rights are granted to create modifications or derivatives of this document.
- No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document, with the sole exception that code included in this document is licensed subject to the Zero-Clause BSD open source license (0BSD), <http://opensource.org/licenses/0BSD>.
- No license (express or implied, by estoppel or otherwise) to any intellectual property rights is granted by this document.
- Code names are used by Intel to identify products, technologies, or services that are in development and not publicly available. These are not "commercial" names and not intended to function as trademarks.
- © Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.